

AgroGenAI: Integrated Real-Time Crop Stress Telemetry and Pathogen Classification System for Precision Agriculture on Bhumi

Technical Project Implementation Report

Author: Sundalam Vijayalaxmi Kalidas

Student ID: 2026CS017 — Department of Computer Engineering

MindForgeAI Division / Chatake Innoworks Pvt. Ltd.

Date: July 2026

Executive Summary

Early detection of crop stress and plant pathogens is critical to preventing widespread agricultural yield loss. Traditional monitoring methods rely on manual field inspection, leading to delayed interventions. This report presents an end-to-end Machine Learning framework and interactive deployment architecture designed to monitor environmental telemetry and classify crop diseases in real time. Using a dataset combining PlantVillage foliage imagery and multi-sensor agronomic variables ($N = 54,303$), we trained and validated a Random Forest model achieving a weighted F_1 -score of 0.88. The platform integrates a real-time Streamlit dashboard featuring a dynamic HTML5 Canvas simulation engine that physically models plant growth, wilting, and discoloration for five major crop types—Cotton, Soybean, Sugarcane, Tomato, and Wheat—based on calculated stress indices.

Keywords: Precision Agriculture, Machine Learning, Random Forest, Stress Index, Streamlit, HTML5 Canvas Simulation, Bhumi Telemetry.

1 Problem Statement & Objectives

Agricultural productivity on *Bhumi* faces persistent threats from unmitigated drought stress and pathogen outbreaks. Late diagnosis of fungal, bacterial, or environmental stress directly compromises food security and farmer livelihoods.

The primary objectives of this project are:

1. Develop an end-to-end machine learning pipeline capable of classifying crop health states and stress levels using combined imagery and agronomic telemetry.
2. Address structural data challenges such as extreme class imbalance and multi-collinearity without introducing data leakage.
3. Deploy an interactive web application that provides real-time "What-If" environmental stress simulation with live, crop-specific animations.

2 System Architecture & Methodology

2.1 Data Processing and Feature Engineering

The underlying dataset integrates 54,303 instances across 12 agronomic parameters. To mitigate multi-collinearity among atmospheric variables, Vapor Pressure Deficit (*VPD*) was derived to replace correlated raw temperature and humidity metrics. Categorical growth stages were ordinally encoded, while foliage spot vectors were reduced using Principal Component Analysis (*PCA*, $k = 5$).

Missing values were handled via median imputation, and numeric scales were normalized using `RobustScaler`. The preprocessing and modeling operations were wrapped inside an isolated pipeline to eliminate cross-validation leakage:

```
1 from sklearn.pipeline import Pipeline
2 from sklearn.impute import SimpleImputer
3 from sklearn.preprocessing import StandardScaler
4 from sklearn.ensemble import RandomForestClassifier
5
6 # Isolated preprocessing and model pipeline
7 pipeline = Pipeline([
8     ('imputer', SimpleImputer(strategy='median')),
9     ('scaler', StandardScaler()),
10    ('clf', RandomForestClassifier(
11        class_weight='balanced',
12        random_state=42
13    ))
14 ])
```

Listing 1: Scikit-Learn ML Pipeline Implementation

2.2 Stress Index Calculation

Environmental stress (S) is continuously evaluated using a weighted linear representation bounded within $[0, 100]$:

$$S = \max(0, \min(100, \lfloor 1.3T + 2.2D - 0.9M \rfloor)) \quad (1)$$

where T represents ambient temperature ($^{\circ}\text{C}$), D denotes consecutive days without rainfall, and M is the soil moisture percentage.

3 Experimental Evaluation & Results

Models were evaluated using 5-Fold Stratified Cross-Validation (**StratifiedKFold**) to preserve class distribution across folds. As missing a diseased state (False Negative) carries a high operational penalty, weighted F_1 -score and Recall served as primary metrics.

Table 1: Model Performance Comparison Across Candidates

Model Candidate	Precision	Recall	Weighted F_1	Execution Speed
Dummy Classifier	0.12	0.20	0.18	< 1 ms
Logistic Regression (L_2)	0.74	0.71	0.72	12 ms
XGBoost Classifier	0.86	0.87	0.865	45 ms
Random Forest (Selected)	0.89	0.88	0.88	8 ms

The selected Random Forest classifier achieved a sealed test set F_1 -score of 0.88, demonstrating robust generalization across diverse target categories without overfitting.

4 Interactive Dashboard & Simulation Architecture

4.1 Field Telemetry Dashboard

The user interface was constructed using Streamlit to allow real-time parameter tweaking. Users can select crop types (Cotton, Soybean, Sugarcane, Tomato, Wheat) and modify telemetry sliders to observe immediate status updates.

4.2 Dynamic Crop Simulation Engine

To replace static indicators, an interactive HTML5 Canvas element was embedded directly into the Streamlit rendering pipeline using `streamlit.components.v1.html`. The renderer uses a dynamic JavaScript request animation loop to physically alter plant geometry based on the stress index S :






Figure 1: Streamlit field telemetry controls and real-time risk assessment overview.

- **Optimal Zone** ($S < 35$): Deep green foliage (#4CAF50), tall stem height (160px), active sway animation.
- **Moderate Stress** ($35 \leq S < 70$): Stunted stem (120px), yellowing leaves (#D4AC0D), slight curvature toward *Bhumi*.
- **Critical Danger** ($S \geq 70$): Severe shrinkage (75px), foliage browning (#795548), and drooping ($\theta = 1.2$ rad).



(a) Optimal Growth State ($S < 35$)



(b) Critical Drought Stress ($S \geq 70$)

Figure 2: Interactive HTML5 Canvas simulator dynamic visual response under varying stress indices on *Bhumi*.

5 System Integration & Deployment Architecture

To bridge technical Machine Learning inference with practical agronomic decision-making on *Bhumi*, the platform follows a lightweight, edge-compatible deployment model.

5.1 Interactive Streamlit State & Canvas Bridge

The web application utilizes Streamlit’s reactive session state architecture to maintain real-time bidirectional parameter flow between UI widgets and the visualization engine:

- **Input Processing:** Telemetry sliders dynamically recalculate the mathematical Stress Index (S) on state change without requiring full page refreshes.
- **HTML5 Canvas Injection:** The physical plant renderer is embedded using `streamlit.components.v1.html`. Calculated stress metrics (S) and crop state vectors are passed as JSON payloads directly into the JavaScript DOM event loop.
- **Mathematical Renderer Control:** The Canvas uses dynamic trigonometry ($\sin(\theta)$), dynamic stem length scaling, and leaf color interpolation from RGB (76, 175, 80) to (121, 85, 72) to simulate physiological wilting dynamically in real time.

5.2 Micro-Services and Scalability Considerations

The inference pipeline is packaged as a modular, stateless Scikit-Learn pipeline object serialized via `joblib`. This decoupling allows the model core to be exposed via lightweight REST APIs (FastAPI) for micro-controller telemetric ingestion (e.g., ESP32 or Raspberry Pi field nodes) while maintaining sub-10ms response latency.

6 Operational Conclusion, Impact & Future Scope

6.1 Operational Impact

The AgroGenAI platform demonstrates a robust end-to-end framework for early crop disease detection and stress prediction ($F_1 = 0.88$). By translating multidimensional telemetry logs and visual foliage vector data into an actionable, visual stress scale, extension officers and farmers on *Bhumi* can execute targeted micro-irrigation and dynamic fungicide applications prior to irreversible yield degradation.

6.2 Limitations

- **Stationary Calibration:** The current feature distribution relies heavily on stationary telemetry ranges and requires localized threshold recalibration when deployed across significantly different soil textures or geographic micro-climates.
- **Telemetry Noise Sensitivity:** Extreme single-sensor measurement anomalies (e.g., faulty soil moisture probes) can temporarily distort calculated stress indices without multi-sensor consensus validation.

6.3 Future Enhancements

1. **Multimodal Deep Learning Integration:** Future iterations will merge Convolutional Neural Network (CNN) leaf feature extraction directly into the tabular pipeline via late-fusion architectures.
2. **Satellite Remote Sensing Integration:** Incorporating Sentinel-2 multi-spectral satellite imagery to scale stress index calculations from single telemetry points to regional farm grid mappings.
3. **Edge Hardware Deployment:** Quantizing the Random Forest classifier into ONNX format for direct local execution on low-power IoT field gateways without continuous internet connectivity.

References

- [1] Hughes, D., & Salathé, M. (2015). An open access repository of plant pathology images for AI-based disease detection. *arXiv preprint arXiv:1511.08060*.
- [2] Breiman, L. (2001). Random forests. *Machine learning*, 45(1), 5-32.
- [3] Pedregosa, F., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825-2830.